

Practical 4

Aim : To implement a Machine Learning Classification model using a Logistic regression algorithm

Data Preprocessing

```
In [23]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, accuracy_score
from sklearn.metrics import precision_recall_fscore_support
from sklearn import metrics
```

```
In [24]: data = pd.read_csv(r"practical4.csv")
```

```
In [25]: data.head()
```

```
Out[25]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunci
0	6	148	72	35	0	33.6	0.
1	1	85	66	29	0	26.6	0.
2	8	183	64	0	0	23.3	0.
3	1	89	66	23	94	28.1	0.
4	0	137	40	35	168	43.1	2.

```
In [26]: data.isna().sum()
```

```
Out[26]: Pregnancies      0
Glucose      0
BloodPressure  0
SkinThickness  0
Insulin      0
BMI          0
DiabetesPedigreeFunction  0
Age          0
Outcome      0
dtype: int64
```

In [27]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Pregnancies                          768 non-null    int64
1   Glucose                              768 non-null    int64
2   BloodPressure                        768 non-null    int64
3   SkinThickness                        768 non-null    int64
4   Insulin                              768 non-null    int64
5   BMI                                  768 non-null    float64
6   DiabetesPedigreeFunction             768 non-null    float64
7   Age                                  768 non-null    int64
8   Outcome                              768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

In [28]: data.describe()

Out[28]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Diab
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	

Parameter setup

In [29]: x= data.iloc[:, [0,7]].values
y=data.iloc[:, -1].values

In [30]: x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.25,random_

Model Training

```
In [31]: classifier= LogisticRegression()
classifier.fit(x_train, y_train)
```

Out[31]: LogisticRegression()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

Making Predication

```
In [32]: y_pred= classifier.predict(x_test)
```

```
In [33]: print(y_test)
```

```
[1 0 0 1 0 0 1 1 0 0 1 1 0 0 0 0 1 0 0 0 1 1 0 0 0 0 0 0 0 0 0 0 0 0 1 1
 0 0 0 0 0 0 1 1 0 0 1 1 1 0 0 1 0 0 0 0 1 1 1 1 0 0 1 1 1 1 0 0 0 0 0 0
 1 0 0 0 0 0 0 0 0 0 0 0 1 0 1 0 0 0 0 0 0 1 0 1 1 0 0 0 0 0 1 0 0 0 1 0
 1 1 1 1 1 0 0 0 1 0 0 0 0 0 0 0 1 0 0 0 0 0 1 0 1 0 1 1 0 0 0 0 0 1 0 0 0
 0 1 0 1 0 0 1 0 0 0 1 1 1 1 0 0 0 1 0 0 0 0 0 0 1 1 0 0 0 0 0 0 1 1 0 1 1
 0 1 1 1 0 0 0]
```

```
In [34]: print(y_pred)
```

```
[0 0 0 1 0 0 0 0 0 1 0 1 1 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 1 0 0 1 0 0 0 0 1
 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0
 0 0 0 0 1 0 0 0 1 0 1 0 1 0 0 0 0 1 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 1 0 0 0 0 1 1 0 0 0 0 1 0 0 0 0 0 0 0 0 1 0 1 0 0 0 0 1 0 0 0
 0 0 0 0 0 1 0]
```

Confusion Matrix

```
In [35]: from sklearn.metrics import confusion_matrix
cm= confusion_matrix(y_test,y_pred)
print(cm)
```

```
[[117  13]
 [ 49  13]]
```

```
In [36]: precision_recall_fscore_support(y_test, y_pred, average='macro')
```

Out[36]: (0.6024096385542168, 0.5548387096774193, 0.542997542997543, None)

```
In [37]: precision_recall_fscore_support(y_test, y_pred, average='micro')
```

Out[37]: (0.6770833333333334, 0.6770833333333334, 0.6770833333333334, None)

```
In [38]: precision_recall_fscore_support(y_test, y_pred, average='weighted')
```

```
Out[38]: (0.638679718875502, 0.6770833333333334, 0.6306690212940212, None)
```

```
In [39]: accuracy_score(y_test,y_pred)
```

```
Out[39]: 0.6770833333333334
```

```
In [40]: Accuracy = metrics.accuracy_score(y_test,y_pred)
Accuracy
```


```
Out[40]: 0.6770833333333334
```

```
In [41]: classifier.intercept_
classifier.coef_
```

```
Out[41]: array([[0.05180623, 0.03442028]])
```

Plotting of Confusion Matrix

```
In [42]: import seaborn as sns
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cbar=False)
plt.title('Confusion Matrix')
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.show()
```



```
In [43]: from sklearn.metrics import precision_recall_fscore_support
precision, recall, f1_score, _ = precision_recall_fscore_support(y_test, y_p)
print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1_score)
```

```
Precision: [0.70481928 0.5          ]
Recall: [0.9          0.20967742]
F1 Score: [0.79054054 0.29545455]
```